

▼ UCI ML Repo - credit card defaults

Preprocessing

Platform: Python 3, colab.research.google.com

```
1 import datetime
2 import matplotlib.pyplot as plt
3 import pandas as pd
4 import seaborn as sns
5 from google.colab import drive
```

```
1 sns.set(style='whitegrid', context='notebook')
```

▼ Load data

```
1 drive.mount('/content/gdrive', force_remount=False)
```

☞ Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive"

```
1 data = pd.read_csv("/content/gdrive/My Drive/Colab Notebooks/uci-credit-card-defaults/data/defaults.csv", header=
2 data.shape
```

☞ (30000, 25)

```
1 data.head(5)
```

☞

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	...	BILL_AMT4	BILL_AMT5	BILL_AMT6
0	1	20000	2	2	1	24	2	2	-1	-1	...	0	0	0
1	2	120000	2	2	2	26	-1	2	0	0	...	3272	3455	3529
2	3	90000	2	2	2	34	0	0	0	0	...	14331	14948	15564
3	4	50000	2	2	1	37	0	0	0	0	...	28314	28959	29604
4	5	50000	1	2	1	57	-1	0	-1	0	...	20940	19146	19751

```
1 data.tail(5)
```



	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	...	BILL_AMT4	BILL_AMT5	BILL_AMT6
29995	29996	220000	1	3	1	39	0	0	0	0	...	88004	31237	31842
29996	29997	150000	1	3	2	43	-1	-1	-1	-1	...	8979	5190	5250
29997	29998	30000	1	2	2	37	4	3	2	-1	...	20878	20582	21187
29998	29999	80000	1	3	1	41	1	-1	0	0	...	52774	11855	12460
29999	30000	50000	1	2	1	46	0	0	0	0	...	36535	32428	33033

5 rows × 25 columns

▼ Exploration

```
1 data.dtypes
```



ID	int64
LIMIT_BAL	int64
SEX	int64
EDUCATION	int64
MARRIAGE	int64
AGE	int64
PAY_0	int64
PAY_2	int64
PAY_3	int64
PAY_4	int64
PAY_5	int64
PAY_6	int64
BILL_AMT1	int64
BILL_AMT2	int64
BILL_AMT3	int64
BILL_AMT4	int64
BILL_AMT5	int64
BILL_AMT6	int64
PAY_AMT1	int64
PAY_AMT2	int64
PAY_AMT3	int64
PAY_AMT4	int64
PAY_AMT5	int64
PAY_AMT6	int64
default payment next month	int64

```
1 data.describe()
```



	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000


```

1 # explore distribution of numerical features
2 def chart_distribution(df: pd.DataFrame):
3     column_distributions = df.describe()
4     idx = 0
5     for i in column_distributions:
6         fig = plt.figure()
7         ax11 = fig.add_subplot(211)
8         ax21 = fig.add_subplot(212)
9         df.loc[:,i].plot(kind="hist", bins=20, title=i, ax=ax11)
10        column_distributions.loc[["min", "25%", "50%", "75%", "max", "mean"], i].plot(
11            kind="bar", ax=ax21)
12        print(idx, i)
13        plt.show()
14        idx += 1
15 chart_distribution(data)

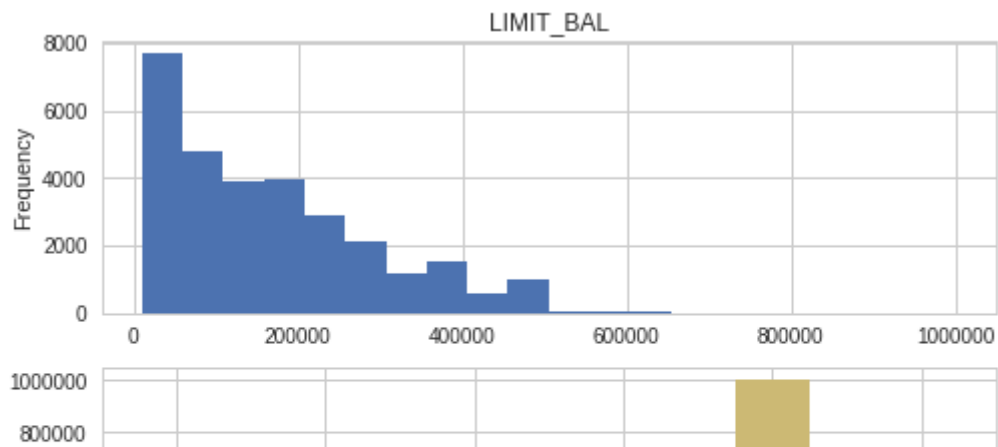
```



0 ID



1 LIMIT_BAL



```
1 # explore distribution of categorical features
2 def chart_categorical(df: pd.DataFrame):
3     idx = 0
4     for col in df:
5         if df[col].dtype == "object":
6             df[col].value_counts().plot(kind="bar")
7             print(idx, col)
```

```

8         plt.show()
9         idx += 1
10 chart_categorical(data)

```

▼ Potential data errors to review:

- education should only have classes 1-4, dataset includes classes 0-6
- marriage should only have classes 1-3, dataset includes classes 0-3
- PAY_1 - PAY_6: represent payment delay in months. -1 is duly paid. now includes -2 and 0 as well
- BILL_AMT1 - BILL_AMT6: includes negative numbers and outliers
- PAY_AMT1 - PAY_AMT6: includes outliers

```

1 # rename columns
2 data_clean = data.copy(deep=True)
3 cols = list(data_clean.columns)
4 cols = [i.lower() for i in cols]
5 cols[-1] = "default"
6 data_clean.columns = cols
7 data_clean.columns

```

```

↳ Index(['id', 'limit_bal', 'sex', 'education', 'marriage', 'age', 'pay_0',
        'pay_2', 'pay_3', 'pay_4', 'pay_5', 'pay_6', 'bill_amt1', 'bill_amt2',
        'bill_amt3', 'bill_amt4', 'bill_amt5', 'bill_amt6', 'pay_amt1',
        'pay_amt2', 'pay_amt3', 'pay_amt4', 'pay_amt5', 'pay_amt6', 'default'],
        dtype='object')

```

```

1 # fix education classes
2 data_clean.loc[data_clean["education"]==0, "education"] = 4
3 data_clean.loc[data_clean["education"]==5, "education"] = 4
4 data_clean.loc[data_clean["education"]==6, "education"] = 4
5 data_clean.loc[:, "education"].unique()

```

```

↳ array([2, 1, 3, 4])

```

```

1 # fix marriage classes
2 data_clean.loc[data_clean["marriage"]==0, "marriage"] = 3
3 data_clean.loc[:, "marriage"].unique().

```

```

↳ array([1, 2, 3])

```

▼ Duplicates

```
1 # verify no duplicates exist
2 data_clean.loc[data_clean.duplicated()==True, ]
```

```
↳      id limit_bal sex education marriage age pay_0 pay_2 pay_3 pay_4 ... bill_amt4 bill_amt5 bill_amt6
0 rows x 25 columns
```

▼ Missing data

```
1 def get_missing(df):
2     exploration = pd.DataFrame(df.isnull().sum(),
3                               columns=["no_missing"])
4     exploration = exploration.merge(
5         pd.DataFrame(df.dtypes, columns=["type"]),
6         left_index=True, right_index=True)
7     return exploration.loc[exploration["no_missing"]>0].sort_values(
8         by="no_missing", ascending=False)
9 missing = get_missing(data_clean)
10 print("Column \t # missing values")
11 for i, row in missing.iterrows():
12     print("{} \t {}".format(i, row["no_missing"]))
```

```
↳ Column    # missing values
```

▼ Outliers

```
1 bill_cols = ["bill_amt1", "bill_amt2", "bill_amt3", "bill_amt4", "bill_amt5", "bill_amt6"]
2 data_clean.loc[:,bill_cols].describe()
```

```
↳
```

	bill_amt1	bill_amt2	bill_amt3	bill_amt4	bill_amt5	bill_amt6
count	30000.000000	30000.000000	3.000000e+04	30000.000000	30000.000000	30000.000000
mean	51223.330900	49179.075167	4.701315e+04	43262.948967	40311.400967	38871.760400
std	73635.860576	71173.768783	6.934939e+04	64332.856134	60797.155770	59554.107537
min	-165580.000000	-69777.000000	-1.572640e+05	-170000.000000	-81334.000000	-339603.000000
25%	3558.750000	2984.750000	2.666250e+03	2326.750000	1763.000000	1256.000000
50%	22381.500000	21200.000000	2.008850e+04	19052.000000	18104.500000	17071.000000
75%	67091.000000	64006.250000	6.016475e+04	54506.000000	50190.500000	49198.250000

```

1 for i in bill_cols:
2     bill_gt_600 = data_clean.loc[data_clean[i]>600000, [i]]
3     print("col {} len {} where val > 600k. max {}".format(i, bill_gt_600.shape[0], bill_gt_600.max()[0]))

```

```

☞ col bill_amt1 len 10 where val > 600k. max 964511
col bill_amt2 len 6 where val > 600k. max 983931
col bill_amt3 len 6 where val > 600k. max 1664089
col bill_amt4 len 4 where val > 600k. max 891586
col bill_amt5 len 2 where val > 600k. max 927171
col bill_amt6 len 2 where val > 600k. max 961664

```

All high bill data points assumed to be reasonable based on nearby values

```

1 for i in bill_cols:
2     bill_lt = data_clean.loc[data_clean[i]<-50000, [i]]
3     print("col {} len {} where val < 50k. max {}".format(i, bill_lt.shape[0], bill_lt.min()[0]))

```

```

☞ col bill_amt1 len 2 where val < 50k. max -165580
col bill_amt2 len 2 where val < 50k. max -69777
col bill_amt3 len 2 where val < 50k. max -157264
col bill_amt4 len 4 where val < 50k. max -170000
col bill_amt5 len 3 where val < 50k. max -81334
col bill_amt6 len 8 where val < 50k. max -339603

```

All negative bill data points assumed to be reasonable based on nearby values


```

1 pay_cols = ["pay_amt1", "pay_amt2", "pay_amt3", "pay_amt4", "pay_amt5", "pay_amt6"]
2 data_clean.loc[:, pay_cols].describe()

```

	pay_amt1	pay_amt2	pay_amt3	pay_amt4	pay_amt5	pay_amt6
count	30000.000000	3.000000e+04	30000.000000	30000.000000	30000.000000	30000.000000
mean	5663.580500	5.921163e+03	5225.68150	4826.076867	4799.387633	5215.502567
std	16563.280354	2.304087e+04	17606.96147	15666.159744	15278.305679	17777.465775
min	0.000000	0.000000e+00	0.000000	0.000000	0.000000	0.000000
25%	1000.000000	8.330000e+02	390.00000	296.000000	252.500000	117.750000
50%	2100.000000	2.009000e+03	1800.00000	1500.000000	1500.000000	1500.000000
75%	5006.000000	5.000000e+03	4505.00000	4013.250000	4031.500000	4000.000000
max	873552.000000	1.684259e+06	896040.00000	621000.000000	426529.000000	528666.000000

```

1 for i in pay_cols:
2     pay_gt = data_clean.loc[data_clean[i]>400000, [i]]
3     print("col {} len {} where val > 400k. max {}".format(i, pay_gt.shape[0], pay_gt.max()[0]))

```

```

col pay_amt1 len 5 where val > 400k. max 873552
col pay_amt2 len 7 where val > 400k. max 1684259
col pay_amt3 len 5 where val > 400k. max 896040
col pay_amt4 len 5 where val > 400k. max 621000
col pay_amt5 len 2 where val > 400k. max 426529
col pay_amt6 len 5 where val > 400k. max 528666

```

pay_amtx outliers appear valid due to other nearby datapoints

▼ Correlated features

```

1 # show correlation > threshold
2 def show_feature_correlation(df):
3     df_corr = df.corr()
4     high_correlations = pd.DataFrame(columns=["f1", "f2", "corr"])
5     for i, row in df_corr.iterrows():
6         for j in row.index:
7             if i == j:
8                 continue
9             high_correlations.loc[len(high_correlations), :] = [i, j, abs(df_corr.loc[i, j])]
10    high_correlations = high_correlations.sort_values(by="corr", ascending=False)
11    high_correlations = high_correlations.loc[high_correlations.loc[:, "corr"].duplicated(), :].
12    return high_correlations
13 high_correlations = show_feature_correlation(data_clean)
14 print(high_correlations.head(10))

```

```

↳
      f1      f2      corr
324  bill_amt2  bill_amt1  0.951484
400  bill_amt5  bill_amt6  0.946197
375  bill_amt4  bill_amt5  0.940134
349  bill_amt3  bill_amt2  0.928326
350  bill_amt3  bill_amt4  0.923969
423  bill_amt6  bill_amt4  0.900941
326  bill_amt2  bill_amt4  0.892482
301  bill_amt1  bill_amt3  0.892279
398  bill_amt5  bill_amt3  0.88391
372  bill_amt4  bill_amt1  0.860272


```

```

1 show_corr = 0.90
2 correlation_cols = set()
3 for i, row in high_correlations.iterrows():
4     if row["corr"] > show_corr:
5         correlation_cols.add(row["f1"])
6         correlation_cols.add(row["f2"])
7 correlation_cols
8 if len(correlation_cols) > 1:
9     sns.pairplot(data_clean.loc[:, sorted(correlation_cols)], size=2)
10 plt.show()

```

↳





```

1 # remove feature with correlation higher than 'remove_corr'
2 remove_corr = 0.99
3 for i in range(len(data_clean.columns)):
4     for i, row in high_correlations.iterrows():
5         if row["corr"] > remove_corr:
6             data_clean = data_clean.drop(row["f2"], axis=1)
7             high_correlations = show_feature_correlation(data_clean)
8             break
9 print(high_correlations.head(10))
10 print(data_clean.columns)

```

```

↳
      f1      f2      corr
324  bill_amt2  bill_amt1  0.951484
400  bill_amt5  bill_amt6  0.946197
375  bill_amt4  bill_amt5  0.940134
349  bill_amt3  bill_amt2  0.928326
350  bill_amt3  bill_amt4  0.923969
423  bill_amt6  bill_amt4  0.900941
326  bill_amt2  bill_amt4  0.892482
301  bill_amt1  bill_amt3  0.892279
398  bill_amt5  bill_amt3  0.88391
372  bill_amt4  bill_amt1  0.860272
Index(['id', 'limit_bal', 'sex', 'education', 'marriage', 'age', 'pay_0',
      'pay_2', 'pay_3', 'pay_4', 'pay_5', 'pay_6', 'bill_amt1', 'bill_amt2',
      'bill_amt3', 'bill_amt4', 'bill_amt5', 'bill_amt6', 'pay_amt1',
      'pay_amt2', 'pay_amt3', 'pay_amt4', 'pay_amt5', 'pay_amt6', 'default'],
      dtype='object')

```




▼ Skewed / Imbalanced data


```

1 for i in data_clean.columns:
2     data_clean.loc[:, i].plot(kind="hist", bins=20)
3     plt.show()

```

↳





▼ Final dataset

```
1 data_clean.describe()
```

	id	limit_bal	sex	education	marriage	age	pay_0	pay_2
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	15000.500000	167484.322667	1.603733	1.842267	1.557267	35.485500	-0.016700	-0.133767
std	8660.398374	129747.661567	0.489129	0.744494	0.521405	9.217904	1.123802	1.197186
min	1.000000	10000.000000	1.000000	1.000000	1.000000	21.000000	-2.000000	-2.000000
25%	7500.750000	50000.000000	1.000000	1.000000	1.000000	28.000000	-1.000000	-1.000000
50%	15000.500000	140000.000000	2.000000	2.000000	2.000000	34.000000	0.000000	0.000000
75%	22500.250000	240000.000000	2.000000	2.000000	2.000000	41.000000	0.000000	0.000000
max	30000.000000	1000000.000000	2.000000	4.000000	3.000000	79.000000	8.000000	8.000000

8 rows × 25 columns

```
1 data_clean.columns
```



```
Index(['id', 'limit_bal', 'sex', 'education', 'marriage', 'age', 'pay_0', ...])
```

Notes for training

- convert columns: sex, education, marriage to categorical
- power transform (log) columns: limit_bal, bill_amtx, pay_amtx

▼ Save datasets

```
1 loc = "/content/gdrive/My Drive/Colab Notebooks/uci-credit-card-defaults/data/defaults_clean.csv"
2 data_clean.to_csv(loc, index=False)
3 # verify file saved correctly
4 data_clean_load = pd.read_csv(loc)
5 assert data_clean_load.shape == data_clean.shape
```

